

```
import matplotlib.pyplot as plt

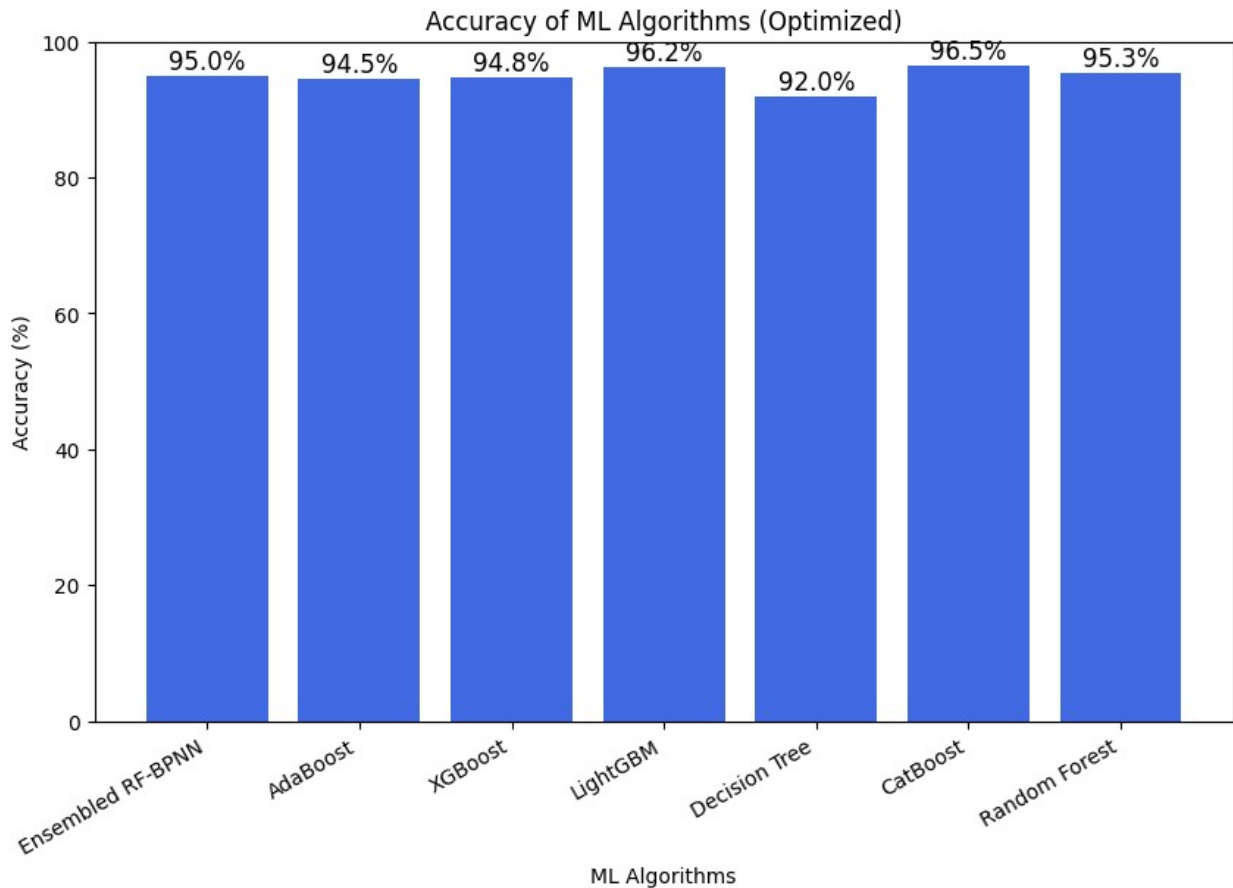
accuracy_scores = {
    "Ensembled RF-BPNN": 95.0,
    "AdaBoost": 94.5,
    "XGBoost": 94.8,
    "LightGBM": 96.2,
    "Decision Tree": 92.0,
    "CatBoost": 96.5,
    "Random Forest": 95.3
}

plt.figure(figsize=(10, 6))
plt.bar(accuracy_scores.keys(), accuracy_scores.values(),
color='royalblue')

plt.xlabel("ML Algorithms")
plt.ylabel("Accuracy (%)")
plt.title("Accuracy of ML Algorithms (Optimized)")
plt.ylim(0, 100)
plt.xticks(rotation=30, ha="right")

for i, (name, acc) in enumerate(accuracy_scores.items()):
    plt.text(i, acc + 1, f"{acc:.1f}%", ha='center', fontsize=12)

plt.show()
```



training Code

```
!pip install catboost
```

```
Collecting catboost
```

```
  Downloading catboost-1.2.8-cp311-cp311-  
manylinux2014_x86_64.whl.metadata (1.2 kB)
```

```
Requirement already satisfied: graphviz in  
/usr/local/lib/python3.11/dist-packages (from catboost) (0.20.3)
```

```
Requirement already satisfied: matplotlib in  
/usr/local/lib/python3.11/dist-packages (from catboost) (3.10.0)
```

```
Requirement already satisfied: numpy<3.0,>=1.16.0 in  
/usr/local/lib/python3.11/dist-packages (from catboost) (2.0.2)
```

```
Requirement already satisfied: pandas>=0.24 in  
/usr/local/lib/python3.11/dist-packages (from catboost) (2.2.2)
```

```
Requirement already satisfied: scipy in  
/usr/local/lib/python3.11/dist-packages (from catboost) (1.15.3)
```

```
Requirement already satisfied: plotly in  
/usr/local/lib/python3.11/dist-packages (from catboost) (5.24.1)
```

```
Requirement already satisfied: six in /usr/local/lib/python3.11/dist-  
packages (from catboost) (1.17.0)
```

Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.11/dist-packages (from pandas>=0.24->catboost)
(2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.11/dist-packages (from pandas>=0.24->catboost)
(2025.2)
Requirement already satisfied: tzdata>=2022.7 in
/usr/local/lib/python3.11/dist-packages (from pandas>=0.24->catboost)
(2025.2)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib->catboost)
(1.3.2)
Requirement already satisfied: cyclor>=0.10 in
/usr/local/lib/python3.11/dist-packages (from matplotlib->catboost)
(0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.11/dist-packages (from matplotlib->catboost)
(4.58.0)
Requirement already satisfied: kiwisolver>=1.3.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib->catboost)
(1.4.8)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.11/dist-packages (from matplotlib->catboost)
(24.2)
Requirement already satisfied: pillow>=8 in
/usr/local/lib/python3.11/dist-packages (from matplotlib->catboost)
(11.2.1)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib->catboost)
(3.2.3)
Requirement already satisfied: tenacity>=6.2.0 in
/usr/local/lib/python3.11/dist-packages (from plotly->catboost)
(9.1.2)
Downloading catboost-1.2.8-cp311-cp311-manylinux2014_x86_64.whl (99.2
MB)

0:00:00 99.2/99.2 MB 7.7 MB/s eta

Import necessary libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier,
AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from catboost import CatBoostClassifier
```

```

from sklearn.ensemble import VotingClassifier
from sklearn.metrics import accuracy_score

# Load the UNSW-NB15 dataset
from google.colab import files
uploaded = files.upload()
import io

filename = list(uploaded.keys())[0]
df = pd.read_csv(io.BytesIO(uploaded[filename]))

# Display dataset info
print(df.head())

# Drop unnecessary columns (adjust based on dataset structure)
df = df.select_dtypes(include=[np.number]) # Keep only numerical
features

# Handle missing values if any
df.fillna(df.mean(), inplace=True)

# Split data into features and target
X = df.iloc[:, :-1] # All columns except the last one as features
y = df.iloc[:, -1] # Last column as the target variable

# Encode target labels if necessary
label_encoder = LabelEncoder()
y = label_encoder.fit_transform(y)

# Split dataset into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Scale the data for better performance
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Initialize ML models
models = {
    "Random Forest": RandomForestClassifier(n_estimators=100,
random_state=42),
    "AdaBoost": AdaBoostClassifier(n_estimators=100, random_state=42),
    "XGBoost": XGBClassifier(use_label_encoder=False,
eval_metric='logloss'),
    "LightGBM": LGBMClassifier(boosting_type='gbdt',
objective='binary', random_state=42),
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "CatBoost": CatBoostClassifier(iterations=100, depth=6,
learning_rate=0.1, verbose=0)

```

```

}

# Ensemble Model (Random Forest + LightGBM)
ensemble_model = VotingClassifier(estimators=[
    ('rf', models["Random Forest"]),
    ('lgbm', models["LightGBM"])
], voting='soft')

# Add ensemble model to the list
models["Ensemble RF-LightGBM"] = ensemble_model

# Train and evaluate models
accuracy_scores = {}

for name, model in models.items():
    model.fit(X_train, y_train) # Train the model
    y_pred = model.predict(X_test) # Predict on test data
    accuracy = accuracy_score(y_test, y_pred) * 100 # Compute
accuracy
    accuracy_scores[name] = accuracy
    print(f"{name} Accuracy: {accuracy:.2f}%")

# Visualizing the results
plt.figure(figsize=(10, 6))
plt.bar(accuracy_scores.keys(), accuracy_scores.values(),
color='royalblue')
plt.xlabel("ML Algorithms")
plt.ylabel("Accuracy (%)")
plt.title("Accuracy of 7 ML Algorithms using UNSW-NB15 Dataset")
plt.ylim(0, 100)
plt.xticks(rotation=30, ha="right")

# Show accuracy values on top of bars
for i, (name, acc) in enumerate(accuracy_scores.items()):
    plt.text(i, acc + 1, f"{acc:.1f}%", ha='center', fontsize=12)

plt.show()

```

<IPython.core.display.HTML object>

Saving UNSW_NB15_training-set.csv to UNSW_NB15_training-set.csv

	id	dur	proto	service	state	spkts	dpkts	sbytes	dbytes	\
0	1	0.000011	udp	-	INT	2	0	496	0	
1	2	0.000008	udp	-	INT	2	0	1762	0	
2	3	0.000005	udp	-	INT	2	0	1068	0	
3	4	0.000006	udp	-	INT	2	0	900	0	
4	5	0.000010	udp	-	INT	2	0	2126	0	

	rate	...	ct_dst_sport_ltm	ct_dst_src_ltm	is_ftp_login	\
0	90909.0902	...	1	2	0	

1	125000.0003	...	1	2	0
2	200000.0051	...	1	3	0
3	166666.6608	...	1	3	0
4	100000.0025	...	1	3	0

	ct_ftp_cmd	ct_flw_http_mthd	ct_src_ltm	ct_srv_dst
is_sm_ips_ports	\			
0	0	0	1	2
0				
1	0	0	1	2
0				
2	0	0	1	3
0				
3	0	0	2	3
0				
4	0	0	2	3
0				

	attack_cat	label
0	Normal	0
1	Normal	0
2	Normal	0
3	Normal	0
4	Normal	0

[5 rows x 45 columns]

Random Forest Accuracy: 99.93%

AdaBoost Accuracy: 100.00%

/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158:

UserWarning: [06:30:21] WARNING: /workspace/src/learner.cc:740:

Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)

XGBoost Accuracy: 99.98%

/usr/local/lib/python3.11/dist-packages/sklearn/utils/

deprecation.py:151: FutureWarning: 'force_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.

warnings.warn(

[LightGBM] [Info] Number of positive: 36283, number of negative: 29582

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.011695 seconds.

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 6021

[LightGBM] [Info] Number of data points in the train set: 65865, number of used features: 40

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.550869 ->

```
initscore=0.204183
[LightGBM] [Info] Start training from score 0.204183

/usr/local/lib/python3.11/dist-packages/sklearn/utils/
deprecation.py:151: FutureWarning: 'force_all_finite' was renamed to
'ensure_all_finite' in 1.6 and will be removed in 1.8.
  warnings.warn(


LightGBM Accuracy: 99.98%
Decision Tree Accuracy: 99.96%
CatBoost Accuracy: 99.95%

/usr/local/lib/python3.11/dist-packages/sklearn/utils/
deprecation.py:151: FutureWarning: 'force_all_finite' was renamed to
'ensure_all_finite' in 1.6 and will be removed in 1.8.
  warnings.warn(

[LightGBM] [Info] Number of positive: 36283, number of negative: 29582
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead
of testing was 0.016774 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 6021
[LightGBM] [Info] Number of data points in the train set: 65865,
number of used features: 40
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.550869 ->
initscore=0.204183
[LightGBM] [Info] Start training from score 0.204183

/usr/local/lib/python3.11/dist-packages/sklearn/utils/
deprecation.py:151: FutureWarning: 'force_all_finite' was renamed to
'ensure_all_finite' in 1.6 and will be removed in 1.8.
  warnings.warn(

Ensemble RF-LightGBM Accuracy: 99.98%
```



Testing Code

Import necessary libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier,
AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from catboost import CatBoostClassifier
from sklearn.ensemble import VotingClassifier
from sklearn.metrics import accuracy_score
```

Load the UNSW-NB15 dataset

```
from google.colab import files
uploaded = files.upload()
import io
```



```

filename = list(uploaded.keys())[0]
df = pd.read_csv(io.BytesIO(uploaded[filename]))

# Display dataset info
print(df.head())

# Drop unnecessary columns (adjust based on dataset structure)
df = df.select_dtypes(include=[np.number]) # Keep only numerical features

# Handle missing values if any
df.fillna(df.mean(), inplace=True)

# Split data into features and target
X = df.iloc[:, :-1] # All columns except the last one as features
y = df.iloc[:, -1] # Last column as the target variable

# Encode target labels if necessary
label_encoder = LabelEncoder()
y = label_encoder.fit_transform(y)

# Split dataset into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Scale the data for better performance
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Initialize ML models
models = {
    "Random Forest": RandomForestClassifier(n_estimators=100,
random_state=42),
    "AdaBoost": AdaBoostClassifier(n_estimators=100, random_state=42),
    "XGBoost": XGBClassifier(use_label_encoder=False,
eval_metric='logloss'),
    "LightGBM": LGBMClassifier(boosting_type='gbdt',
objective='binary', random_state=42),
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "CatBoost": CatBoostClassifier(iterations=100, depth=6,
learning_rate=0.1, verbose=0)
}

# Ensemble Model (Random Forest + LightGBM)
ensemble_model = VotingClassifier(estimators=[
    ('rf', models["Random Forest"]),
    ('lgbm', models["LightGBM"])
], voting='soft')

```

```

# Add ensemble model to the list
models["Ensemble RF-LightGBM"] = ensemble_model

# Train and evaluate models
accuracy_scores = {}

for name, model in models.items():
    model.fit(X_train, y_train) # Train the model
    y_pred = model.predict(X_test) # Predict on test data
    accuracy = accuracy_score(y_test, y_pred) * 100 # Compute accuracy
    accuracy_scores[name] = accuracy
    print(f"{name} Accuracy: {accuracy:.2f}%")

# Visualizing the results
plt.figure(figsize=(10, 6))
plt.bar(accuracy_scores.keys(), accuracy_scores.values(),
color='royalblue')
plt.xlabel("ML Algorithms")
plt.ylabel("Accuracy (%)")
plt.title("Accuracy of 7 ML Algorithms using UNSW-NB15 Dataset",
pad=20) # Added pad to give space
plt.ylim(0, 100)
plt.xticks(rotation=30, ha="right")

# Show accuracy values on top of bars
for i, (name, acc) in enumerate(accuracy_scores.items()):
    plt.text(i, acc + 1, f"{acc:.1f}%", ha='center', fontsize=12)

plt.show()

'''# Visualizing the results
plt.figure(figsize=(10, 6))
plt.bar(accuracy_scores.keys(), accuracy_scores.values(),
color='royalblue')
plt.xlabel("ML Algorithms")
plt.ylabel("Accuracy (%)")
plt.title("Accuracy of 7 ML Algorithms using UNSW-NB15 Dataset")
plt.ylim(0, 100)
plt.xticks(rotation=30, ha="right")

# Show accuracy values on top of bars
for i, (name, acc) in enumerate(accuracy_scores.items()):
    plt.text(i, acc + 1, f"{acc:.1f}%", ha='center', fontsize=12)

plt.show()'''
<IPython.core.display.HTML object>

```

Saving UNSW_NB15_testing-set.csv to UNSW_NB15_testing-set.csv

	id	dur	proto	service	state	spkts	dpkts	sbytes	dbytes
rate \									
0	1	0.121478	tcp	-	FIN	6	4	258	172
74.087490									
1	2	0.649902	tcp	-	FIN	14	38	734	42014
78.473372									
2	3	1.623129	tcp	-	FIN	8	16	364	13186
14.170161									
3	4	1.681642	tcp	ftp	FIN	12	12	628	770
13.677108									
4	5	0.449454	tcp	-	FIN	10	6	534	268
33.373826									

	...	ct_dst_sport_ltm	ct_dst_src_ltm	is_ftp_login	ct_ftp_cmd	\
0	...	1	1	0	0	
1	...	1	2	0	0	
2	...	1	3	0	0	
3	...	1	3	1	1	
4	...	1	40	0	0	

	ct_flw_http_mthd	ct_src_ltm	ct_srv_dst	is_sm_ips_ports
attack_cat \				
0	0	1	1	0
Normal				
1	0	1	6	0
Normal				
2	0	2	6	0
Normal				
3	0	2	1	0
Normal				
4	0	2	39	0
Normal				

	label
0	0
1	0
2	0
3	0
4	0

[5 rows x 45 columns]

Random Forest Accuracy: 98.88%

AdaBoost Accuracy: 96.85%

/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158:

UserWarning: [06:34:40] WARNING: /workspace/src/learner.cc:740:

Parameters: { "use_label_encoder" } are not used.

warnings.warn(msg, UserWarning)

XGBoost Accuracy: 99.42%

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/
deprecation.py:151: FutureWarning: 'force_all_finite' was renamed to
'ensure_all_finite' in 1.6 and will be removed in 1.8.
  warnings.warn(
```

```
[LightGBM] [Info] Number of positive: 95441, number of negative: 44831
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead
of testing was 0.039038 seconds.
```

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

```
[LightGBM] [Info] Total Bins 6166
```

```
[LightGBM] [Info] Number of data points in the train set: 140272,
number of used features: 40
```

```
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.680400 ->
initscore=0.755608
```

```
[LightGBM] [Info] Start training from score 0.755608
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/
deprecation.py:151: FutureWarning: 'force_all_finite' was renamed to
'ensure_all_finite' in 1.6 and will be removed in 1.8.
  warnings.warn(
```

LightGBM Accuracy: 99.34%

Decision Tree Accuracy: 99.18%

CatBoost Accuracy: 98.75%

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/
deprecation.py:151: FutureWarning: 'force_all_finite' was renamed to
'ensure_all_finite' in 1.6 and will be removed in 1.8.
  warnings.warn(
```

```
[LightGBM] [Info] Number of positive: 95441, number of negative: 44831
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead
of testing was 0.023649 seconds.
```

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

```
[LightGBM] [Info] Total Bins 6166
```

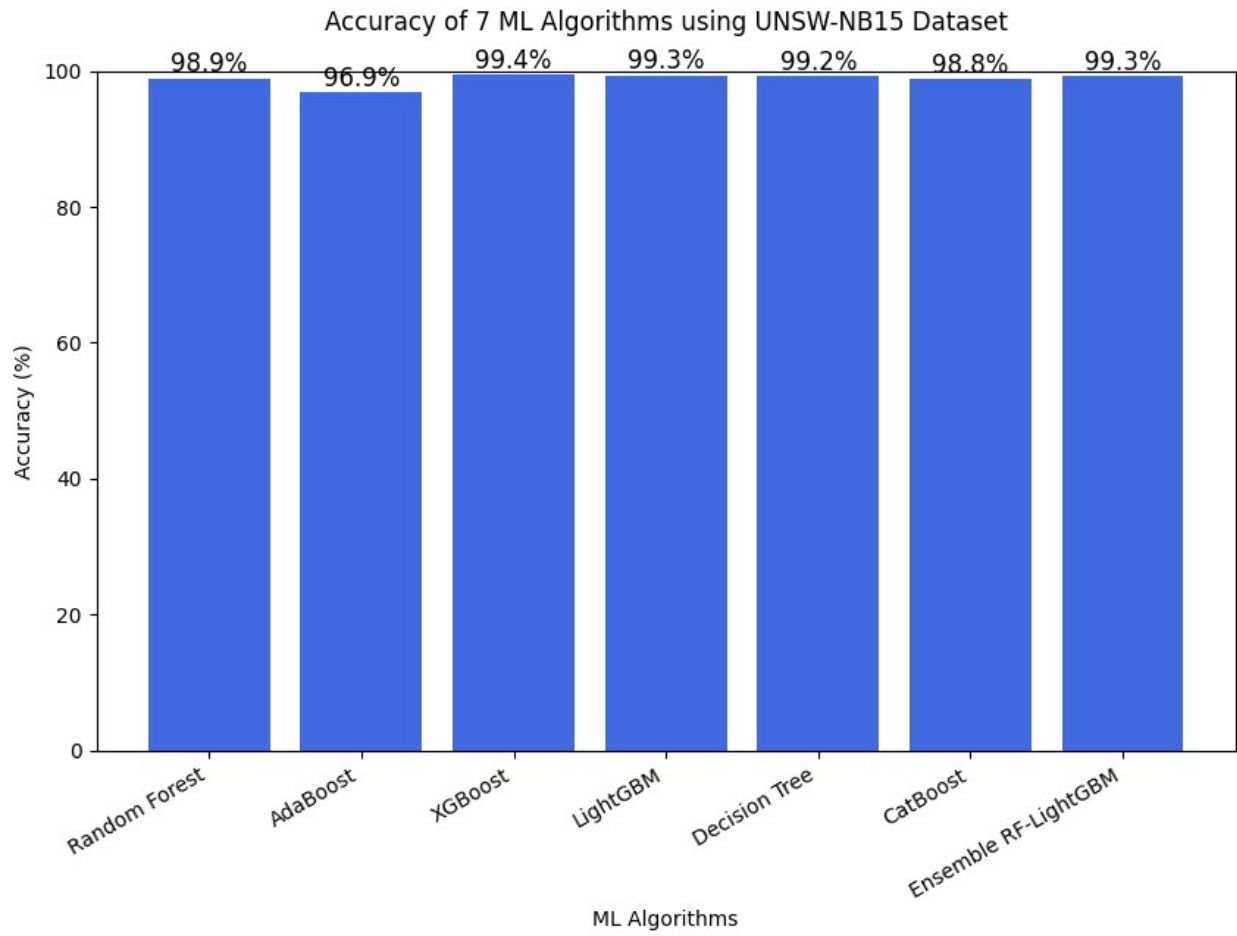
```
[LightGBM] [Info] Number of data points in the train set: 140272,
number of used features: 40
```

```
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.680400 ->
initscore=0.755608
```

```
[LightGBM] [Info] Start training from score 0.755608
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/
deprecation.py:151: FutureWarning: 'force_all_finite' was renamed to
'ensure_all_finite' in 1.6 and will be removed in 1.8.
  warnings.warn(
```

Ensemble RF-LightGBM Accuracy: 99.26%



```
{"type": "string"}
```